Noise Covariance as Representation Geometry: One Metric, Four Architecture Classes

Anonymous Author(s)
Affiliation
Address
email

Abstract

Label-free scientific reconstruction learns from measured observations alone: each observation supplies its own reconstruction supervision, while the instrument-characterized noise covariance determines which residuals are informative. For Gaussian noise, inverse covariance is the likelihood-consistent metric. Rather than proposing a new estimator, we organize classical methods as a pair $(\Sigma^{-1}, \mathcal{F})$ of inverse covariance and representation class, separating measurement geometry from admissible model capacity, and follow this metric through optimal filtering, unrestricted linear projection, Kronecker-separable projection, and factored attention. The linear classes admit an exact geometric characterization. At the separable-projection rung, models sharing an architecture but differing only in the training metric reverse their ranking under raw MSE and the likelihood residual, in a frozen ten-seed experiment and in an independent twenty-seed replication. The reversal persists when the covariance is estimated from $2{,}048$ noise-only calibration events in one fixed sensitivity condition, while the covariance-weighted factors recover the planted channel subspace almost exactly but the temporal subspace only weakly. At the factored-attention rung the reversal is absent across two independent seeds, despite non-degenerate weighting. Noise covariance therefore defines the scientifically relevant geometry, while instrument and representation class determine whether it changes what is learned.

1 Introduction

Many scientific instruments produce abundant measurements but no labels or clean targets: each observation is its own training signal, and an instrument-derived noise model says which discrepancies matter. Uncertainty varies across frequencies, sensors and conditions; mean-squared error (MSE) weighs all residual coordinates equally, yet under colored noise a large residual can lie in a poorly measured direction while a small one is decisive. The measurement covariance is therefore a measurement-model prior fixing the geometry of reconstruction error, distinct from a conservation law or equivariance [Bronstein et al., 2021, Karniadakis et al., 2021].

This complements self-supervised reconstruction without clean targets [Batson and Royer, 2019, Yaman et al., 2020, Gruber et al., 2025]. Augmentation- and masking-based objectives [Chen et al., 2020, He et al., 2022] construct supervision by transforming or withholding part of an observation; here the measured observation supplies the reconstruction target, and an independently calibrated noise model supplies the residual weighting. We ask a controlled question: if an instrument-derived covariance fixes the training metric, how does its effect change as the admissible representation class expands?

Every link is classical, and we state the boundary plainly. Inverse-covariance weighting is generalized least squares [Aitken, 1936] and, at rank one, the matched filter [Turin, 1960, Kay, 1993, Szymkowiak

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

et al., 1993] — which fixes the weighting but assumes a *known*, one-dimensional template. The latent-variable line relaxes the signal model but not the noise model: probabilistic PCA assumes isotropic noise and factor analysis a diagonal covariance, while weighted low-rank models, Mahalanobis low-rank representation, the linear-autoencoder/PCA correspondence and multilinear PCA are all established [Tipping and Bishop, 1999, Rubin and Thayer, 1982, Roweis, 1998, Bailey, 2012, Allen et al., 2014, Zhang et al., 2019, Baldi and Hornik, 1989, Lu et al., 2008]. *We introduce no new estimator.* Our contribution is a precise separation of measurement geometry from representation capacity, with matched tests of that distinction across established linear, multilinear and attention-based classes.

The pair $(\Sigma^{-1}, \mathcal{F})$ makes that separation precise: the measurement likelihood fixes the residual geometry Σ^{-1}, the architecture the admissible class $\mathcal{F}$ — including the axis along which the data are tokenized (Section 3). Holding Σ^{-1} fixed, we walk $\mathcal{F}$ up the four rungs of Section 3. The first three admit an exact geometric characterization (rung 3's joint optimization is nonconvex; its subproblems are exact), including the condition under which the Kronecker restriction attains the unrestricted optimum (Proposition 2); the fourth is examined empirically. Section 4 then measures the metric reversal and its replication, what the weighted factors recover, the sensitivity of the reversal to a covariance estimated from finite calibration data, and the absence of the reversal at rung 4. Nonlinear classes, distribution shift and adaptive sampling remain future work.

2 Covariance-weighted representation geometry

Let $x_i = s_i + n_i \in \mathbb{R}^d$ be mean-subtracted with $n_i \sim \mathcal{N}(0, \Sigma)$, $\Sigma \succ 0$ (mean handling in Appendix B.2). For a candidate $\hat{s}_i$ the Gaussian negative log-likelihood is, up to constants,

$$\chi^2(\hat{s}_i) = (x_i - \hat{s}_i)^\top \Sigma^{-1} (x_i - \hat{s}_i). \tag{1}$$

Whitening, $\tilde{x}_i = \Sigma^{-1/2} x_i$, maps this form to $\|\tilde{x}_i - \hat{\tilde{s}}_i\|_2^2$: Euclidean projection in whitened coordinates, Σ^{-1}-orthogonal projection in measurement coordinates. We report the per-coordinate normalized residual $\bar{\chi}^2 = (Nd)^{-1} \sum_{i=1}^N \chi^2(\hat{s}_i)$, or the per-trace mean where stated.

Why label-free learning works: at the population level, for independent zero-mean signal and noise, $\mathrm{Cov}(\Sigma^{-1/2} x) = \mathrm{Cov}(\Sigma^{-1/2} s) + I$ (Appendix A); adding I shifts eigenvalues, not eigenvectors, so the leading whitened eigenspace of the *noisy* data is the signal's own; at finite N the sample eigenvectors are perturbed. In practice Σ comes from calibration or noise-only records, estimated without evaluation data; a misspecified Σ is a misspecified geometry.

Proposition 1 (Noise-weighted projection bridge)**.** *Under the model above, OF, rank-k covariance-weighted PCA (CW-PCA), and a tied noise-aware linear autoencoder minimize Eq.* (1) *over different admissible subspaces. OF fixes a rank-one template; CW-PCA learns a rank-k subspace; and the tied autoencoder parameterizes the same learned projection when its encoder is the Σ^{-1}-adjoint of its decoder.*

Proof sketch. Whitening reduces the problem to ordinary rank-k approximation, so every minimizer spans a leading eigenspace of the whitened sample covariance, and the tied Σ^{-1}-adjoint encoder follows from the same geometry. Projection onto a *fixed* rank-one template yields the OF amplitude; rank-one CW-PCA instead *learns* the leading whitened direction, and the two coincide under the separable-template and eigengap assumptions of Corollary 1. Details in Appendix A.

Metric reversal. Raw MSE replaces Σ^{-1} by I and is likelihood-correct only when $\Sigma = \sigma^2 I$; under colored noise MSE follows raw variance while Eq. (1) emphasizes high-information directions. Consequently, two equal-rank models can satisfy $\mathrm{MSE}(A) < \mathrm{MSE}(B)$ but $\chi^2(A) > \chi^2(B)$: each is optimal in a different geometry.

3 The architecture ladder

Holding the metric fixed, we walk $\mathcal{F}$ up four rungs: rungs 1–3 in general two-axis form, evidenced synthetically; rung 4 on the short-time Fourier transform (STFT) band–frame plane of a public benchmark, where Σ^{-1} is diagonal in frequency by measurement (Section 4).

Rung 1: rank-one fixed template (optimal filtering). Fixing the signal to a known template s_0, the likelihood-optimal reconstruction is the optimal filter (OF) amplitude [Gatti and Manfredi, 1986, Turin, 1960]; the representation class is a single direction and the metric question has a closed-form answer.

Rung 2: unrestricted rank-k linear. Replacing the fixed template with a learned rank-k subspace gives covariance-weighted PCA (CW-PCA), which is also the tied noise-aware linear autoencoder: both minimize the likelihood residual over the full Grassmannian $\mathrm{Gr}(k, n_1 n_2)$ of a two-axis measurement (Proposition 1), with reconstruction $\hat{s} = PP^\top \Sigma^{-1} x$ for a Σ^{-1}-orthonormal basis P.

Rung 3: Kronecker-separable two-axis projection (NFPA). For axis-aligned noise, $\Sigma = \Sigma_2 \otimes \Sigma_1$, NFPA restricts rung 2's projection to separable factors of a two-axis measurement $Z_i \in \mathbb{R}^{n_1 \times n_2}$ — sensor channel $\times$ time in the controlled synthetic study, frequency band $\times$ time frame on the public benchmark. On the per-axis whitened observation $\tilde{Z}_i = \Sigma_1^{-1/2} Z_i \Sigma_2^{-1/2}$ it learns orthonormal bases $U_1 \in \mathbb{R}^{n_1 \times k_1}$ and $U_2 \in \mathbb{R}^{n_2 \times k_2}$, with code $C_i = U_1^\top \tilde{Z}_i U_2$ and reconstruction $\hat{\tilde{Z}}_i = U_1 U_1^\top \tilde{Z}_i U_2 U_2^\top$ mapped back through $\Sigma^{1/2}$; the joint representation is $\mathrm{range}(U_1) \otimes \mathrm{range}(U_2)$, and the admissible family is the embedded product $\mathrm{Gr}(k_1, n_1) \times \mathrm{Gr}(k_2, n_2) \hookrightarrow \mathrm{Gr}(k_1 k_2, n_1 n_2)$.

Proposition 2 (Kronecker-restricted bridge)**.** *NFPA minimizes the CW-PCA/tied-autoencoder objective restricted to Kronecker-separable two-axis subspaces. Its optimum equals the unrestricted CW-PCA optimum if and only if at least one rank-$k_1 k_2$ CW-PCA minimizer has such a separable subspace.*

Proof sketch. In whitened coordinates the factored map is the orthogonal projector onto $\mathrm{range}(U_1) \otimes \mathrm{range}(U_2)$ (vec identity, Appendix A), so NFPA is the whitened rank-$k_1 k_2$ objective restricted to that product; equality holds precisely when that set contains an unrestricted minimizer, and otherwise the gap is the cost of the architecture prior, not a failure of the metric — a cost Appendix B.2 measures by fitting rungs 2 and 3 at matched rank. This is a covariance-aware member of the multilinear-PCA family [Lu et al., 2008], not a new factorization. Even a full-rank factor leaves the restriction active — with $k_1 = n_1$ the family is $\mathbb{R}^{n_1} \otimes \mathrm{range}(U_2)$, still separable (Appendix B.2 shows this case empirically); at $k_1 = k_2 = 1$ rung 1 is recovered under Corollary 1's assumptions (Appendix A).

Rung 4: factored attention. Rung 3's restriction on the band–frame plane reads as $\mathrm{range}(U_f) \otimes \mathrm{range}(U_m)$; rung 4's transformer performs intra-band temporal attention over frames, then cross-band attention. The attention model is an *analogous axis-factored architecture* in the same coordinates (no containment of the Kronecker family is claimed), and Σ^{-1} is diagonal in exactly the basis the first factor indexes: on this benchmark a stationary-PSD model fits the readout noise ($\chi^2/\mathrm{dof} = 1.0031$, all gates passed; Section 4). The tokenization principle: choose the token axis in which Σ^{-1} is diagonal, and the likelihood metric becomes a per-token weight rather than a dense coupling; accordingly, real and imaginary components are two features of each of 328 physical-frequency tokens, not 656 unrelated rows. The rung-4 model is a transformer [Vaswani et al., 2017] over frame patches [Nie et al., 2023], with rotary temporal embeddings [Su et al., 2024] and learned band embeddings, following a design demonstrated in prior work on calorimetric event reconstruction [Sjostrom, 2026]; the architecture is a vehicle, not a contribution, and the metric question is posed empirically (Section 4).

4 Evidence

NFPA under channel compression. Rung 3 is consolidated at $C = 8 > k_c = 2$ channels, matched rank $k = k_c k_t = 6$: NFPA, an isotropically trained factored baseline of the same architecture, CW-PCA, and PCA on the same seeds. NFPA and the isotropic baseline differ only in the training metric, isolating it, and it reverses: in the frozen ten-seed experiment NFPA is better under the likelihood residual (-0.0192, 95% interval $[-0.0194, -0.0190]$) and worse under raw MSE ($+0.0034$, $[+0.0034, +0.0034]$), both excluding zero. An independent twenty-seed replication of the identical configuration agrees in sign and magnitude (-0.0190, $[-0.0192, -0.0189]$; $+0.0034$); cohorts are reported separately, any pooled figure being post-hoc rather than a thirty-seed confirmatory experiment. The twenty-seed white-noise run is an identity check: at $\Sigma = \sigma^2 I$ the objectives coincide

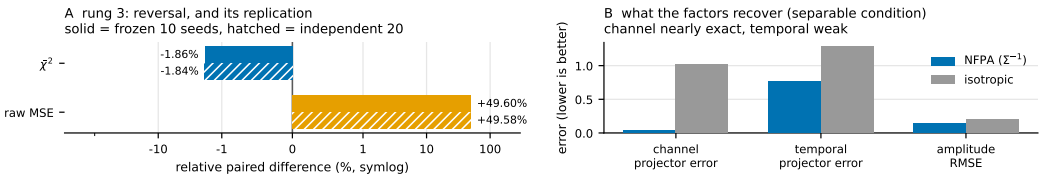


Figure 1: **A**, rung 3: the reversal and its replication. Relative paired differences, NFPA minus the isotropic baseline; solid = frozen ten seeds, hatched = an independent twenty-seed replication. The metrics point in *opposite* directions — the reversal — and the cohorts agree. Symmetric-log axis; cohorts are never pooled. **B**, what the factors recover at the separable condition: the channel subspace almost exactly, the temporal one only weakly; its three groups carry different units, so bars are comparable within a group only. Generated from archived summaries.

up to scale, and its exact zero shows no spurious contrast. Estimating the covariance from $2{,}048$ noise-only calibration events preserves the reversal in all ten seeds (-0.0191, $[-0.0193, -0.0189]$; $+0.0034$), degrading NFPA's likelihood residual and raw MSE by 0.0078% and 0.0345% of its own oracle held-out values — consistent with invariance to a global whitening scale and an accurately estimated covariance *shape*, but one fixed condition, not a robustness sweep (Table 3).

Rung 4 (exploratory): the metric in a factored-attention network. One factored-attention architecture is trained on the TIDMAD readout under two objectives differing only in the metric — `T-mse` on the unwhitened spectrogram, `T-chi2` on the Σ^{-1}-weighted residual — at matched data, budget and evaluation grid, at two independent seeds. *The reversal does not appear at either seed.* The Σ^{-1}-trained arm is better on *both* criteria in both runs: held-out paired differences of -1.15 and -1.17, both $\times 10^{-4}$, on the likelihood residual and -57.2 and -57.4 on raw MSE, the latter opposite to prediction; all four signs agree across seeds, which bounds run-to-run variation rather than characterising it. The release's coarse score does not separate the arms (-1.0×10^{-3}, intervals spanning zero), and both sit far below the published baselines (Appendix B.1). A degenerate metric would explain it, and we tested for one: J spans a factor 441 and the whitening weights have coefficient of variation 0.69, the same order as the rung-3 geometry where the reversal *is* observed (1.10; white control 0), so the absence is a property of this instantiation, not of the ladder.

What the covariance-weighted factors recover. The metric also changes what the representation *is*. At the separable condition NFPA recovers the planted channel subspace almost exactly (0.035 projector error against the isotropic 1.021) while temporal recovery stays weak (0.768 against 1.280) — an asymmetry a joint number would hide. All seven paired contrasts favour NFPA, every across-seed bootstrap interval excludes zero and all ten seeds agree in sign, with no multiplicity correction. Two are delayed-predeclared — the channel and temporal projector errors — and five are post-hoc: two joint projector errors and the decoding of amplitude, rise and decay time. No recovery threshold was predeclared either way, so these are continuous estimates, not a verdict (Appendix B.2).

Transfer of the noise model. On a public gravitational-wave strain release [LIGO Scientific Collaboration and Virgo Collaboration, 2016, Macleod et al., 2021], a whitened score matching its reference form to 10^{-12} and calibrated on synthetic stationary noise (dispersion 1.15) is about eight times as dispersed on real H1 and L1 strain (9.42, 8.83; global-PSD arm 1.77, 7.29). A covariance model can be correctly implemented, validated synthetically, and still fail to transfer: implementation identity is not calibration.

Limitations. Rung 4 rests on two seeds per arm, and its band-limited arms are not comparable to the published full-band baselines; with the objectives there measurably non-degenerate, the absence remains unexplained — capacity, optimization trajectory and implicit regularization all remain candidates. Semantic recovery is asymmetric and five of its seven contrasts are post-hoc; the linear-rung TIDMAD reversal is not replicated at a confirmatory file; stationary-PSD adequacy is per-instrument (adequate on the TIDMAD readout, $\chi^2/\mathrm{dof} = 1.0031$, failing on GW strain); and the primary experiment's oracle covariance is narrowed, not removed, by one estimated-covariance condition — no sweep over calibration sizes, covariance families or misspecification.

References

Alexander C. Aitken. On least squares and linear combination of observations. *Proceedings of the Royal Society of Edinburgh*, 55:42–48, 1936.

Genevera I. Allen, Logan Grosenick, and Jonathan Taylor. A generalized least-square matrix decomposition. *Journal of the American Statistical Association*, 109(505):145–159, 2014. doi: 10.1080/01621459.2013.852978.

Stephen Bailey. Principal component analysis with noisy and/or missing data. *Publications of the Astronomical Society of the Pacific*, 124(919):1015–1023, 2012.

Pierre Baldi and Kurt Hornik. Neural networks and principal component analysis: Learning from examples without local minima. *Neural Networks*, 2(1):53–58, 1989.

Joshua Batson and Loic Royer. Noise2Self: Blind denoising by self-supervision. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 524–533, 2019. URL `https://proceedings.mlr.press/v97/batson19a.html`.

Michael M. Bronstein, Joan Bruna, Taco Cohen, and Petar Veličković. Geometric deep learning: Grids, groups, graphs, geodesics, and gauges. *arXiv preprint arXiv:2104.13478v2*, 2021. URL `https://arxiv.org/abs/2104.13478`.

Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020. URL `https://arxiv.org/abs/2002.05709`.

Emilio Gatti and Paolo F. Manfredi. Processing the signals from solid-state detectors in elementary-particle physics. *La Rivista del Nuovo Cimento*, 9(1):1–146, 1986. doi: 10.1007/BF02822156.

Nadja Gruber, Johannes Schwab, Markus Haltmeier, Ander Biguri, Clemens Dlaska, and Gyeongha Hwang. Noisier2Inverse: Self-supervised learning for image reconstruction with correlated noise. *arXiv preprint arXiv:2503.19468*, 2025. doi: 10.48550/arXiv.2503.19468. URL `https://arxiv.org/abs/2503.19468`.

Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16000–16009, 2022. URL `https://arxiv.org/abs/2111.06377`.

George Em Karniadakis, Ioannis G. Kevrekidis, Lu Lu, Paris Perdikaris, Sifan Wang, and Liu Yang. Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440, 2021. doi: 10.1038/s42254-021-00314-5.

Steven M. Kay. *Fundamentals of Statistical Signal Processing, Volume I: Estimation Theory*. Prentice Hall, 1993.

Tamara G. Kolda and Brett W. Bader. Tensor decompositions and applications. *SIAM Review*, 51(3):455–500, 2009. doi: 10.1137/07070111X.

LIGO Scientific Collaboration and Virgo Collaboration. Data release for event GW150914. Gravitational Wave Open Science Center, 2016. URL `https://gwosc.org/events/GW150914/`. Accessed 12 July 2026.

Haiping Lu, Konstantinos N. Plataniotis, and Anastasios N. Venetsanopoulos. MPCA: Multilinear principal component analysis of tensor objects. *IEEE Transactions on Neural Networks*, 19(1):18–39, 2008. doi: 10.1109/TNN.2007.901277.

Duncan M. Macleod, Joseph S. Areeda, Scott B. Coughlin, Thomas J. Massinger, and Alexander L. Urban. GWpy: A Python package for gravitational-wave astrophysics. *SoftwareX*, 13:100657, 2021. doi: 10.1016/j.softx.2021.100657.

Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *International Conference on Learning Representations (ICLR)*, 2023. URL https://arxiv.org/abs/2211.14730.

Sam Roweis. EM algorithms for PCA and SPCA. *Advances in Neural Information Processing Systems*, 10:626–632, 1998.

Donald B. Rubin and Dorothy T. Thayer. EM algorithms for ML factor analysis. *Psychometrika*, 47 (1):69–76, 1982.

Sebastian Sjostrom. Adapting transformers for light dark matter event reconstruction. MSci project report, Imperial College London, 2026. Unpublished; available on request.

Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024. doi: 10.1016/j.neucom.2023.127063.

A. E. Szymkowiak, R. L. Kelley, S. H. Moseley, and C. K. Stahle. Signal processing for microcalorimeters. *Journal of Low Temperature Physics*, 93(3–4):281–285, 1993. doi: 10.1007/BF00693433.

Michael E. Tipping and Christopher M. Bishop. Probabilistic principal component analysis. *Journal of the Royal Statistical Society: Series B*, 61(3):611–622, 1999.

George L. Turin. An introduction to matched filters. *IRE Transactions on Information Theory*, 6(3): 311–329, 1960.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems 30*, pages 5998–6008, 2017. URL https://arxiv.org/abs/1706.03762.

Burhaneddin Yaman, Seyed Amir Hossein Hosseini, Steen Moeller, Jutta Ellermann, Kamil Uğurbil, and Mehmet Akçakaya. Self-supervised learning of physics-guided reconstruction neural networks without fully sampled reference data. *Magnetic Resonance in Medicine*, 84(6):3172–3191, 2020. doi: 10.1002/mrm.28378.

Chihao Zhang, Kuo Gai, and Shihua Zhang. Matrix normal PCA for interpretable dimension reduction and graphical noise modeling. *arXiv preprint arXiv:1911.10796*, 2019. doi: 10.48550/arXiv.1911.10796.

A Proof details

A.1 Noise-weighted projection bridge

Let $W = \Sigma^{-1/2}P$. The constraint $P^\top \Sigma^{-1} P = I_k$ becomes $W^\top W = I_k$, and the reconstruction satisfies

$$\Sigma^{-1/2}\hat{s} = \Sigma^{-1/2} P P^\top \Sigma^{-1} x = W W^\top \tilde{x}.$$

The empirical loss is therefore $N^{-1}\sum_i \|\tilde{x}_i - WW^\top \tilde{x}_i\|_2^2$. By Rayleigh–Ritz, every minimizer spans a leading rank-k eigenspace of the whitened sample covariance. Mapping back yields the CW-PCA subspace. The objective depends on W only through $WW^\top$, so W is identified up to an orthogonal rotation.

A.2 Whitened signal-plus-identity decomposition

For $x = s + n$ with $s \perp n$, both zero mean, and $\mathrm{Cov}(n) = \Sigma$: independence gives $\mathrm{Cov}(x) = \mathrm{Cov}(s) + \Sigma$, and conjugating by $\Sigma^{-1/2}$,

$$\mathrm{Cov}(\Sigma^{-1/2}x) = \mathrm{Cov}(\Sigma^{-1/2}s) + I.$$

Adding I shifts every eigenvalue by one and changes no eigenvector, so the eigenvectors of the whitened data covariance are those of the whitened *signal* covariance, in the same order. Rank-k CW-PCA on noisy data therefore recovers the leading whitened signal eigenspace at the population level, which is the precise sense in which the likelihood metric supplies a self-supervised training signal.

A.3 Kronecker-restricted bridge

For every $Z \in \mathbb{R}^{C\times T}$, the vectorization identity gives

$$\mathrm{vec}(U_c U_c^\top Z U_t U_t^\top) = (U_t U_t^\top \otimes U_c U_c^\top)\mathrm{vec}(Z).$$

The Kronecker product is a symmetric idempotent matrix of rank $k_c k_t$, hence the orthogonal projector onto $\mathrm{range}(U_t)\otimes\mathrm{range}(U_c)$. NFPA is thus the ordinary whitened rank-$k_c k_t$ projection loss restricted to product subspaces. The restricted and unrestricted minima coincide exactly when an unrestricted minimizing projector belongs to this subset.

A.4 Rank-one correspondence

For $Z_i = A_i v t^\top + \epsilon_i$ with isotropic whitened noise and a unique leading signal eigendirection separated by a nonzero eigengap, the population second moment has leading vectorized direction $(t \otimes v)/(\|t\|\|v\|)$, which is feasible for rank-one NFPA. Its projection coefficient is proportional to $\langle \mathrm{vec}(Z_i), t \otimes v\rangle$, the whitened matched-filter coefficient and hence the OF amplitude up to normalization.

Corollary 1 (Rank-one correspondence)**.** *Assume isotropic whitened noise and a unique leading population signal eigendirection separated by a nonzero eigengap. If that direction is generated by a separable template $\tilde{s} = vt^\top$, population NFPA with $k_1 = k_2 = 1$ learns $v/\|v\|$ and $t/\|t\|$ up to signs; its scalar code is proportional to the OF amplitude of Section 2.*

B Experiment details and reporting limits

The controlled NFPA experiment uses $C = 2$ channels, $T = 256$ samples, a sampling rate of $1\,\mathrm{kHz}$, channel weights $(1.0, 0.65)$, and amplitude $A \sim \mathcal{N}(1, 0.3^2)$. The noise is Brownian with $S_n(f) \propto 1/f^2$ and is whitened using the matched temporal frequency-domain filter. The train and test sets contain 800 and 300 events. NFPA and IsoAE use $(k_c, k_t) = (2, 3)$, while CW-PCA and IsoPCA use rank six.

Because $C = k_c = 2$, $U_c U_c^\top = I_2$ in this experiment: the reconstruction restricts a temporal subspace shared by both channels but does not compress the channel axis. Principal angles between NFPA and CW-PCA are $0.0°$, $59.3°$, $64.6°$, $73.8°$, $84.8°$, and $87.8°$ (mean $61.7°$). The corresponding

Table 1: Implementation check: the single-split $C = k_c = 2$ run reported by the earlier draft, retained for provenance. $\bar{\chi}^2$ is the per-coordinate normalized whitened residual $\bar{\chi}^2$ of Section 2; $\bar{\theta}$ is the mean principal angle to CW-PCA in whitened space.

Method	Training	Class	$\bar{\chi}^2$	Raw MSE	$\bar{\theta}$
CW-PCA	Σ^{-1}	full	0.9929	4.95×10^{-3}	0°
NFPA	Σ^{-1}	factored	0.9916	4.95×10^{-3}	61.7°
IsoAE	MSE	factored	0.9995	4.58×10^{-9}	86.0°
IsoPCA	MSE	full	1.0007	4.33×10^{-9}	—

Table 2: TIDMAD coarse denoising scores (higher is better) from published baselines, the archived linear estimate, and this paper's two rung-4 arms, which share architecture, data, and budget and differ only in the training objective. The two arms return bit-identical values because the upstream aggregation rounds the raw score to 0.01 before the logarithm and both band-limited arms fall in the same bin; rescoring without the rounding gives 0.0241 vs. 0.0231 in band, a contrast whose bootstrap interval spans zero (Section 4). Band-limited arms are not comparable to the full-band published rows.

Method	Coarse score
FCNet	6.55
Transformer	4.18
PUNet	3.84
none	1.10 (1.1026)
linear Wiener (archived)	-1.8025
`T-mse`	-2.7708
`T-chi2`	-2.7708

IsoAE–CW-PCA angles are 82.1°, 83.7°, 86.0°, 86.6°, 88.2°, and 89.4° (mean 86.0°). These are point estimates from one seed and split, produced by the companion single-split demonstration — a different experiment and code path from the multi-seed cohorts, and therefore *not* the quantities corrected by E11 (Table 3); the E11 values must not be read as replacements for them. The exact held-out raw-MSE values for IsoAE and IsoPCA are re-exported at full precision from the archived run (4.58×10^{-9} and 4.33×10^{-9}).

B.1 TIDMAD rung-4 baselines

A submission-strength empirical study should repeat this comparison across seeds and train/test splits and sweep SNR, noise color, channel count, rank pair, and controlled nonseparability. It should report confidence intervals for held-out $\bar{\chi}^2$, raw MSE, and subspace angles. No such uncertainty estimate is claimed here.

B.2 Predeclared confirmatory protocol

Mean handling, all experiments: the training mean is estimated once, applied unchanged to held-out observations, and added back after reconstruction.

The multi-seed confirmatory study predeclared three hypotheses. **Executed:** H1 (metric reversal at fixed architecture): under colored or correlated noise, NFPA and an isotropically trained multilinear projection with the same (k_c, k_t) should disagree in ranking, with paired across-seed confidence intervals excluding zero in the predicted directions. H2 (white-noise identity check): when $\Sigma = \sigma^2 I$ the two objectives coincide up to scale, so covariance-aware and Euclidean projectors must agree to numerical tolerance; this validates the pipeline rather than the hypothesis. **Executed later (delayed):** H3 (active restriction): with $C > k_c$ and $T > k_t$, NFPA should recover the planted channel and temporal subspaces for separable signals. This was not run with the original blocks; it was executed afterwards on the archived fits, each of which reproduces its archived held-out residual to $\leq 10^{-9}$ relative, so the analysis is computed on the archived fits rather than on refits that merely resemble them. Because the protocol set no numerical recovery threshold, the result is reported as continuous projector errors with bootstrap intervals that resample seeds and no binary verdict.

Truth-referenced parameter decoding (amplitude, rise time, decay time) was proposed in a block dropped before execution; it is reported at the archived $\alpha = 0.1$ condition as *post-hoc exploratory*, and at the separable $\alpha = 0$ condition alongside the delayed H3 analysis. An earlier design specified 30 seeds; the executed, frozen protocol is 10 seeds per the 2026-08-08 amendment, fixed before any confirmatory result existed. The primary configuration uses $C = 8$ channels, $T = 256$ samples, ranks $(k_c, k_t) = (2, 3)$, $800/400/2000$ training/validation/test events, 10 independent seeds, 10 ALS restarts per seed, SNR 10 under $\mathrm{SNR}^2 = \mathbb{E}[s^\top \Sigma^{-1} s]$, Brownian $1/f^2$ temporal noise with channel correlation, oracle covariance, and nonseparability $\alpha = 0.1$. Results enter the main text only if every acceptance gate holds: all seeds complete without post-result configuration changes; the fixed-architecture reversal intervals exclude zero in the predicted directions; the white-noise check passes its predeclared equivalence margin; and the table and figure are generated directly from archived seed-level results.

Rung 3 against the unrestricted rung-2 optimum. At the frozen $C = 8 > k_c = 2$ configuration NFPA's held-out likelihood residual is lower than unrestricted CW-PCA's by 0.0026 ($[0.0025, 0.0026]$). This held-out finite-sample difference does not imply that a restricted training optimum beats the unrestricted one: Proposition 2 compares the fitted objective over nested classes and gives the condition for equal optima.

Semantic recovery, per method. At the separable condition NFPA attains amplitude RMSE 0.143 against the isotropic baseline's 0.198 (decoded-amplitude correlation 0.70 against 0.13), and lower rise- and decay-time errors. The same analysis at the archived $\alpha = 0.1$ condition is post-hoc exploratory and agrees in direction on all seven contrasts.

The factored representation (Section 3, generic form) has $O(k_1 n_1 + k_2 n_2)$ basis parameters rather than $O(k_1 k_2 n_1 n_2)$ for an unstructured basis, trading approximation flexibility for a two-axis inductive bias related to classical tensor factorizations [Kolda and Bader, 2009, Lu et al., 2008]. Alternating spectral updates solve each factor subproblem exactly, although the joint product-manifold objective remains nonconvex.

B.3 Round 10: estimated covariance, and the corrected angle diagnostic

Table 3: Round 10. **E10b** is a single frozen sensitivity condition: the separable covariance is estimated from $2{,}048$ independent noise-only calibration events rather than taken as oracle, at the otherwise unchanged $C = 8 > k_c = 2$ configuration. The reversal survives in all ten seeds. The degradation percentages are relative to NFPA's *own* corresponding archived oracle held-out value. This is one condition, not a sweep across calibration sizes, covariance families or misspecification. **E11** is a post-hoc diagnostic correction of the whitened-basis vectorization used for principal angles; only the `_v2` quantities carry scientific meaning.

Quantity	Estimate
E10b — estimated separable covariance, $N_{\mathrm{cal}} = 2048$, 10 seeds	
$\Delta\bar{\chi}^2$, NFPA − isotropic	-0.019122 $[-0.019340, -0.018914]$
ΔMSE, NFPA − isotropic	$+0.0034178$ $[+0.0033963, +0.0034393]$
NFPA $\bar{\chi}^2$ degradation vs. oracle	7.90×10^{-5} (0.0078%)
NFPA MSE degradation vs. oracle	3.55×10^{-6} (0.0345%)
Product-scale ratio vs. oracle	0.9991
Channel / temporal shape error	0.0282 / 0.0197
Kronecker relative Frobenius error	0.0365
E11 — corrected principal angles to CW-PCA (degrees)	
Frozen, $n = 10$	NFPA 62.89, isotropic 86.71, difference -23.82
Replication, $n = 20$	NFPA 63.03, isotropic 86.73, difference -23.71

E11 note: the correction leaves NFPA unchanged and shifts the isotropic baseline by $0.81°$ in the frozen cohort and $0.73°$ in the replication cohort, so no archived scientific conclusion depended on the defect. Displayed intervals are 95% bootstrap intervals with seed as the resampling unit. Cohorts are never pooled for confirmatory reading.

E10b repeats the frozen $C = 8 > k_c = 2$ configuration with one change: the separable covariance is estimated from $N_{\mathrm{cal}} = 2048$ independent noise-only calibration events instead of being supplied as

an oracle. Estimation proceeds by sample estimation of each factor, the regularization already used for the oracle path, and a symmetric product-scale calibration $\hat{v} = \sqrt{v_c v_t}$ with each factor divided by $\sqrt{\hat{v}}$; that correction was frozen before any full result existed. The reversal survives in all ten seeds, and the degradation relative to NFPA's own oracle held-out values is small but resolved — consistent with the fit being invariant to a global whitening scale (the estimated product scale is within 0.1% of oracle) while the covariance *shape* is estimated accurately. This is a single condition: it does not bound behaviour across calibration sizes, covariance families or deliberate misspecification, and no such sweep is claimed.

E11 corrects a vectorization defect in the whitened-basis helper used for principal angles, which unstacked basis columns in the wrong order and so mixed channel and temporal indices. The corrected quantities carry the `_v2` suffix. The archival implementation and the legacy columns of the earlier experiments are deliberately left unchanged, so that previously archived tables continue to mean what they meant when written; new analysis uses the corrected path only. The practical consequence of the defect is small: NFPA's mean angle is identical to the legacy value, and the isotropic baseline moves by under one degree in both cohorts, so no archived conclusion depended on it. E11 is a post-hoc diagnostic correction and contributes no confirmatory evidence; in particular it is not a second rung-4 seed — that role belongs to E7.

Rung-2 numbers carried from the companion work. In the 30-seed rank-two benchmark (512 traces of dimension 256 per seed), PCA and CW-PCA had held-out weighted residuals 48.419 and 47.869 (paired difference 0.550, $[0.530, 0.571]$) while raw MSE preferred the opposite reconstruction, their whitened subspaces 86.25° apart. The ten-seed white-noise control gave a maximum angle of 1.48×10^{-6} degrees. On the TIDMAD file the amplitude RMSE fell from 0.1209 to 0.0880.